

# PampaR v9: Brain-Inspired Territorial Architecture

A Language Model with Interpretable Routing


Lucas Ricardo Mella Chillemi · Independent Researcher · January 2026

## 💡 Abstract

PampaR is a language model architecture inspired by the functional organization of the human brain. Unlike standard transformers that treat all computations uniformly, PampaR introduces **Territorial Processing** where specialized neural modules are organized into functional territories coordinated by a central **Thalamus** that routes tokens using a hybrid approach: 70% explicit rules (LLAVES) and 30% learned attention.

**Key Result:** ~45 perplexity on WikiText-103 with only 14M parameters, outperforming LSTM (24M, PPL 69.1) and approaching Transformer-XL Small (24M, PPL 54.5) with 42% fewer parameters.

## ✂ Architecture Overview



## 🧠 Territories

PampaR organizes computation into 4 functional territories, each containing specialized modules:

| Territory | Modules | Function | Brain Analogy |
|-----------|---------|----------|---------------|
|-----------|---------|----------|---------------|

|                    |                       |                                     |                                |
|--------------------|-----------------------|-------------------------------------|--------------------------------|
| <b>Expresivo</b>   | Language, Creativity  | Fluent text generation, novel ideas | Broca's area, Right hemisphere |
| <b>Contextual</b>  | Context               | Working memory, coherence           | Hippocampus, Prefrontal cortex |
| <b>Formal</b>      | Logic                 | Deductive reasoning, rules          | Prefrontal cortex              |
| <b>Estructural</b> | Patterns, Mathematics | Sequences, numbers, structure       | Parietal lobe                  |

 **LLAVES: Interpretable Routing**

LLAVES (Spanish for "keys") are explicit, interpretable routing rules based on token patterns:

| Module      | Activation Patterns                              |
|-------------|--------------------------------------------------|
| Language    | "the", "a", "of", "that", articles, prepositions |
| Mathematics | digits 0-9, "+", "-", "=", "%", operators        |
| Logic       | "if", "then", "because", "therefore"             |
| Patterns    | repetitions, sequences, regular structures       |
| Context     | pronouns, references, temporal markers           |
| Creativity  | adjectives, metaphors, novel combinations        |

The routing formula combines rules with learning:

```
routing_weights = λ × LLAVES_activation + (1 - λ) × learned_attention where λ = 0.7
```

**⇒ Bidirectional Frontiers**

Territories communicate via 6 bidirectional frontier connections with learned gates:

| Connection               | Weight            | Function                 |
|--------------------------|-------------------|--------------------------|
| Expresivo ↔ Contextual   | 0.8 (High)        | Narrative needs context  |
| Expresivo ↔ Formal       | 0.5 (Medium)      | Structured argumentation |
| Expresivo ↔ Estructural  | 0.4 (Low)         | Creative structure       |
| Contextual ↔ Formal      | 0.6 (Medium-high) | Logical context          |
| Contextual ↔ Estructural | 0.5 (Medium)      | Patterns in context      |
| Formal ↔ Estructural     | 0.7 (High)        | Mathematical logic       |

 **Model Configuration**

| Parameter          | Value      | Training      | Value          |
|--------------------|------------|---------------|----------------|
| Total Parameters   | 14M        | Dataset       | WikiText-103   |
| Hidden Dimension   | 256        | Hardware      | GTX 1650 (4GB) |
| Attention Heads    | 8          | Optimizer     | AdamW          |
| Territorial Blocks | 6          | Learning Rate | 3e-4           |
| Context Length     | 512 tokens | Precision     | FP16 mixed     |
| Vocabulary         | 32K (BPE)  | Training Time | ~8 hours       |

## 📊 Results Comparison

| Model                | Parameters | PPL ↓      | Year |
|----------------------|------------|------------|------|
| LSTM (Merity et al.) | 24M        | 69.1       | 2018 |
| Transformer-XL Small | 24M        | 54.5       | 2019 |
| <b>PampaR v9</b>     | <b>14M</b> | <b>~45</b> | 2026 |
| GPT-2 Small          | 125M       | 35.1       | 2019 |

## 📋 Key Advantages

- **Interpretability:** LLAVES provide explicit routing rationale
- **Parameter Efficiency:** Competitive PPL with 42% fewer parameters
- **Modularity:** Easy to add or modify specialized territories
- **Biological Plausibility:** Reflects brain's functional organization

## 📋 Limitations & Future Work

- Scale: Not yet tested at 1B+ parameters
- Tasks: Evaluated mainly on language modeling; needs downstream evaluation
- LLAVES Design: Currently manual; could benefit from automatic rule learning

## 🔧 Reproducibility

**GitHub:** [github.com/lucasmella-stack/PAMPAr-o1](https://github.com/lucasmella-stack/PAMPAr-o1)

**HuggingFace:** [huggingface.co/lucasmella-stack/PAMPAr-o1](https://huggingface.co/lucasmella-stack/PAMPAr-o1)

**License:** AGPL-3.0-or-later · **Copyright:** © 2024-2026 Lucas Ricardo Mella Chillemi

